

ECGR 5106 Homework 5: Transformers

Lauren Bourque (Student ID: 801264851)

April 26, 2024

GitHub Link

[Click here to view the code](#)

Problem 1.

Sequences of 10:

The model structure can be seen in Figure 1 below, providing the computational complexity such as number of parameters for each layer:

Figure 1: Model Structure

Layer (type:depth-idx)	Output Shape	Param #
CharTransformer	[1900, 44]	--
└─Embedding: 1-1	[1900, 10, 128]	5,632
└─TransformerEncoder: 1-2	[1900, 10, 128]	--
└─ModuleList: 2-1	--	--
└─TransformerEncoderLayer: 3-1	[1900, 10, 128]	593,024
└─TransformerEncoderLayer: 3-2	[1900, 10, 128]	593,024
└─TransformerEncoderLayer: 3-3	[1900, 10, 128]	593,024
└─Linear: 1-3	[1900, 44]	5,676
Total params: 1,790,380		
Trainable params: 1,790,380		
Non-trainable params: 0		
Total mult-adds (G): 3.03		
Input size (MB): 0.15		
Forward/backward pass size (MB): 1129.12		
Params size (MB): 6.37		
Estimated Total Size (MB): 1135.64		

The training loss, validation loss, validation accuracy, and total execution time can be seen in Figure 2 below:

Figure 2: Training Loss + Accuracy

```

Epoch 5, Loss: 2.829401969909668, Validation Loss: 2.6700103282928467, Validation Accuracy: 0.27941176295280457, Execution Time: 18.270447731018066 seconds
Epoch 10, Loss: 2.552777051925659, Validation Loss: 2.523683786392212, Validation Accuracy: 0.23529411852359772, Execution Time: 16.170519828796387 seconds
Epoch 15, Loss: 2.440364122390747, Validation Loss: 2.4505257606506348, Validation Accuracy: 0.27941176295280457, Execution Time: 16.71851110458374 seconds
Epoch 20, Loss: 2.3740100860595703, Validation Loss: 2.412954807281494, Validation Accuracy: 0.2710084021091461, Execution Time: 16.37756371498108 seconds
Epoch 25, Loss: 2.321685791015625, Validation Loss: 2.368964910507202, Validation Accuracy: 0.29201680421829224, Execution Time: 17.226407289505005 seconds
Epoch 30, Loss: 2.2885217666625977, Validation Loss: 2.3425052165985107, Validation Accuracy: 0.2899159789085388, Execution Time: 16.2433762550354 seconds
Epoch 35, Loss: 2.2667012214660645, Validation Loss: 2.328082799911499, Validation Accuracy: 0.2836134433746338, Execution Time: 16.507776260375977 seconds
Epoch 40, Loss: 2.246110200881958, Validation Loss: 2.3244121074676514, Validation Accuracy: 0.29201680421829224, Execution Time: 16.902719259262085 seconds
Epoch 45, Loss: 2.23183012008667, Validation Loss: 2.316819667816162, Validation Accuracy: 0.3025210201740265, Execution Time: 16.58337688446045 seconds
Epoch 50, Loss: 2.2328975200653076, Validation Loss: 2.319289445877075, Validation Accuracy: 0.28151261806488037, Execution Time: 17.38715934753418 seconds
Total Execution Time: 834.1697237491608 seconds

```

Sequences of 20:

The model structure can be seen in Figure 3 below, providing the computational complexity such as number of parameters for each layer:

Figure 3: Model Structure

```

=====
Layer (type:depth-idx)                               Output Shape                                Param #
=====
CharTransformer                                       [1892, 44]                                --
├─Embedding: 1-1                                     [1892, 20, 128]                           5,632
├─TransformerEncoder: 1-2                             [1892, 20, 128]                           --
│   └─ModuleList: 2-1
│       └─TransformerEncoderLayer: 3-1               [1892, 20, 128]                           593,024
│           └─TransformerEncoderLayer: 3-2           [1892, 20, 128]                           593,024
│               └─TransformerEncoderLayer: 3-3       [1892, 20, 128]                           593,024
└─Linear: 1-3                                         [1892, 44]                                5,676
=====
Total params: 1,790,380
Trainable params: 1,790,380
Non-trainable params: 0
Total mult-adds (G): 3.01
=====
Input size (MB): 0.30
Forward/backward pass size (MB): 2248.06
Params size (MB): 6.37
Estimated Total Size (MB): 2254.73
=====

```

The training loss, validation loss, validation accuracy, and total execution time can be seen in Figure 4 below:

Figure 4: Training Loss + Accuracy

```
Epoch 5, Loss: 2.902303695678711, Validation Loss: 2.8105857372283936, Validation Accuracy: 0.20464135706424713, Execution Time: 33.94730591773987 seconds
Epoch 10, Loss: 2.629087448120117, Validation Loss: 2.5793728828430176, Validation Accuracy: 0.2405063360929489, Execution Time: 33.82742261886597 seconds
Epoch 15, Loss: 2.49463152885437, Validation Loss: 2.50187349319458, Validation Accuracy: 0.24261602759361267, Execution Time: 33.6160569190979 seconds
Epoch 20, Loss: 2.4157893657684326, Validation Loss: 2.4567251205444336, Validation Accuracy: 0.27426159381866455, Execution Time: 33.36484980583191 seconds
Epoch 25, Loss: 2.3571343421936035, Validation Loss: 2.4352853298187256, Validation Accuracy: 0.25738397240638733, Execution Time: 33.260478019714355 seconds
Epoch 30, Loss: 2.320802688598633, Validation Loss: 2.3991174697875977, Validation Accuracy: 0.2637130916118622, Execution Time: 33.400601863861084 seconds
Epoch 35, Loss: 2.2873716354370117, Validation Loss: 2.372976541519165, Validation Accuracy: 0.27848100662231445, Execution Time: 32.386797189712524 seconds
Epoch 40, Loss: 2.270092725753784, Validation Loss: 2.3724288940429688, Validation Accuracy: 0.2552742660045624, Execution Time: 32.659749269485474 seconds
Epoch 45, Loss: 2.2507073879241943, Validation Loss: 2.361520290374756, Validation Accuracy: 0.2805907130241394, Execution Time: 33.32268929481506 seconds
Epoch 50, Loss: 2.2436301708221436, Validation Loss: 2.3631558418273926, Validation Accuracy: 0.2763713002204895, Execution Time: 32.71213507652283 seconds
Total Execution Time: 1658.8146982192993 seconds
```

Sequences of 30:

The model structure can be seen in Figure 5 below, providing the computational complexity such as number of parameters for each layer:

Figure 5: Model Structure

```
=====
Layer (type:depth-idx)                Output Shape                Param #
=====
CharTransformer                        [1884, 44]                  --
├─Embedding: 1-1                      [1884, 30, 128]            5,632
├─TransformerEncoder: 1-2             [1884, 30, 128]            --
│   └─ModuleList: 2-1                --
│       └─TransformerEncoderLayer: 3-1 [1884, 30, 128]            593,024
│           └─TransformerEncoderLayer: 3-2 [1884, 30, 128]            593,024
└─Linear: 1-3                        [1884, 44]                  5,676
=====
Total params: 1,197,356
Trainable params: 1,197,356
Non-trainable params: 0
Total mult-adds (G): 2.01
=====
Input size (MB): 0.45
Forward/backward pass size (MB): 2257.85
Params size (MB): 4.26
Estimated Total Size (MB): 2262.56
=====
```

The training loss, validation loss, validation accuracy, and total execution time can be seen in Figure 6 below:

Figure 6: Training Loss + Accuracy

```
Epoch 5, Loss: 2.7724125385284424, Validation Loss: 2.739386796951294, Validation Accuracy: 0.22669491171836853, Execution Time: 34.22837591171265 seconds
Epoch 10, Loss: 2.542022705078125, Validation Loss: 2.5789334774017334, Validation Accuracy: 0.23516948521137238, Execution Time: 34.438804388046265 seconds
Epoch 15, Loss: 2.441145420074463, Validation Loss: 2.5189452171325684, Validation Accuracy: 0.22669491171836853, Execution Time: 34.14780068397522 seconds
Epoch 20, Loss: 2.3785905838012695, Validation Loss: 2.4940552711486816, Validation Accuracy: 0.21822033822536469, Execution Time: 33.687227964401245 seconds
Epoch 25, Loss: 2.3335747718811035, Validation Loss: 2.459984302520752, Validation Accuracy: 0.24788135290145874, Execution Time: 34.31617975234985 seconds
Epoch 30, Loss: 2.3007843494415283, Validation Loss: 2.446418523788452, Validation Accuracy: 0.24788135290145874, Execution Time: 33.46302628517151 seconds
Epoch 35, Loss: 2.271259307861328, Validation Loss: 2.438570499420166, Validation Accuracy: 0.25, Execution Time: 34.39435172080994 seconds
Epoch 40, Loss: 2.24881911277771, Validation Loss: 2.4344990253448486, Validation Accuracy: 0.24788135290145874, Execution Time: 34.640878200531006 seconds
Epoch 45, Loss: 2.238251209259033, Validation Loss: 2.4346768856048584, Validation Accuracy: 0.24364407360553741, Execution Time: 34.25514793395996 seconds
Epoch 50, Loss: 2.218820810317993, Validation Loss: 2.438627004623413, Validation Accuracy: 0.25211864709854126, Execution Time: 33.24434781074524 seconds
Total Execution Time: 1713.871531009674 seconds
```

Overall, the sequences of 10 provided the best validation accuracy.

Problem 2.

Sequences of 20:

The training loss, validation loss, validation accuracy, and total execution time can be seen in Figure 7 below:

Figure 7: Training Loss + Accuracy

```
Epoch 1, Train Loss: 2.511201670412319, Validation Loss: 2.488522491706801, Validation Accuracy: 0.2612519724573232
Epoch 2, Train Loss: 2.482445878528348, Validation Loss: 2.4822470452655185, Validation Accuracy: 0.2686128245588868
Epoch 3, Train Loss: 2.4770272818857997, Validation Loss: 2.478883504696546, Validation Accuracy: 0.2658334528762014
Epoch 4, Train Loss: 2.4738775046240087, Validation Loss: 2.4746318428130936, Validation Accuracy: 0.2690924903170277
Epoch 5, Train Loss: 2.471594839454501, Validation Loss: 2.478012214208364, Validation Accuracy: 0.2635830583847368
Epoch 6, Train Loss: 2.4701393740263686, Validation Loss: 2.472646136346121, Validation Accuracy: 0.26820936737914214
Epoch 7, Train Loss: 2.4691859109383576, Validation Loss: 2.4724686447097586, Validation Accuracy: 0.26885489886673364
Epoch 8, Train Loss: 2.468294793710256, Validation Loss: 2.474320502965903, Validation Accuracy: 0.26847385597475254
Epoch 9, Train Loss: 2.4675045568175884, Validation Loss: 2.471005383394247, Validation Accuracy: 0.2692673217615837
Epoch 10, Train Loss: 2.46693564340282, Validation Loss: 2.4720575176751214, Validation Accuracy: 0.2689580045904461
Epoch 11, Train Loss: 2.466498386396464, Validation Loss: 2.4715782235845607, Validation Accuracy: 0.2688593817242863
Epoch 12, Train Loss: 2.465959999794222, Validation Loss: 2.471158846227361, Validation Accuracy: 0.2685679959833596
Epoch 13, Train Loss: 2.46548342567967, Validation Loss: 2.46966001108821, Validation Accuracy: 0.26867110170707215
Epoch 14, Train Loss: 2.4651649310862105, Validation Loss: 2.4685175512223823, Validation Accuracy: 0.2689804188782097
Epoch 15, Train Loss: 2.464841829073376, Validation Loss: 2.4692465988126893, Validation Accuracy: 0.26932559890976904
Epoch 16, Train Loss: 2.4646516742673463, Validation Loss: 2.4690229703254465, Validation Accuracy: 0.27023113613541816
Epoch 17, Train Loss: 2.4643821579160625, Validation Loss: 2.468950020601718, Validation Accuracy: 0.2684648902596471
Epoch 18, Train Loss: 2.4640671391515165, Validation Loss: 2.4699773849330797, Validation Accuracy: 0.268451441686989
Epoch 19, Train Loss: 2.4636682573443522, Validation Loss: 2.468704862516625, Validation Accuracy: 0.26930766747955814
Epoch 20, Train Loss: 2.4635618337342016, Validation Loss: 2.4685923818036157, Validation Accuracy: 0.26942422177592884

Results for sequence length: 20

Transformer Model:
Training Loss: 2.4635618337342016
Validation Loss: 2.4685923818036157
Validation Accuracy: 0.26942422177592884
Execution Time: 1314.9288251399994 seconds
```

Sequences of 30:

The training loss, validation loss, validation accuracy, and total execution time can be seen in Figure 8 below:

Figure 8: Training Loss + Accuracy

```
Epoch 1, Train Loss: 2.512194059487107, Validation Loss: 2.4756848642289833, Validation Accuracy: 0.2676000143454014
Epoch 2, Train Loss: 2.4838817431735323, Validation Loss: 2.4719237631709325, Validation Accuracy: 0.26871178295407677
Epoch 3, Train Loss: 2.478510127512279, Validation Loss: 2.4683206546966496, Validation Accuracy: 0.2675775996557104
Epoch 4, Train Loss: 2.4751962940781014, Validation Loss: 2.4687589308614744, Validation Accuracy: 0.2681738304014919
Epoch 5, Train Loss: 2.473546355335893, Validation Loss: 2.4682709120912323, Validation Accuracy: 0.26773450248354763
Epoch 6, Train Loss: 2.4718424746413343, Validation Loss: 2.4671372768473514, Validation Accuracy: 0.2689135151612961
Epoch 7, Train Loss: 2.470636532797343, Validation Loss: 2.464689446360919, Validation Accuracy: 0.26831728441551456
Epoch 8, Train Loss: 2.4699963924009096, Validation Loss: 2.4620981746349484, Validation Accuracy: 0.2706035827640002
Epoch 9, Train Loss: 2.4690966822511187, Validation Loss: 2.464325696504086, Validation Accuracy: 0.2700880449011064
Epoch 10, Train Loss: 2.468264890605929, Validation Loss: 2.465640468090537, Validation Accuracy: 0.27017322072193234
Epoch 11, Train Loss: 2.4679549866790444, Validation Loss: 2.4628491697050268, Validation Accuracy: 0.2688059246507791
Epoch 12, Train Loss: 2.4673739529063012, Validation Loss: 2.4621495505061777, Validation Accuracy: 0.27048702637760685
Epoch 13, Train Loss: 2.466970974379752, Validation Loss: 2.4636309491381776, Validation Accuracy: 0.2695456094105833
Epoch 14, Train Loss: 2.4667180668880824, Validation Loss: 2.462107895770196, Validation Accuracy: 0.2680841716427278
Epoch 15, Train Loss: 2.4664919807405283, Validation Loss: 2.4625976337561326, Validation Accuracy: 0.26871626589201497
Epoch 16, Train Loss: 2.466095617554658, Validation Loss: 2.4621888012276996, Validation Accuracy: 0.26898524216830744
Epoch 17, Train Loss: 2.4659133389559305, Validation Loss: 2.4618011659907815, Validation Accuracy: 0.27005666433553893
Epoch 18, Train Loss: 2.4654697663647176, Validation Loss: 2.461008305828967, Validation Accuracy: 0.2690121397959367
Epoch 19, Train Loss: 2.4652400905483542, Validation Loss: 2.460980305145065, Validation Accuracy: 0.26958595585202716
Epoch 20, Train Loss: 2.4649602108898487, Validation Loss: 2.460609679763612, Validation Accuracy: 0.2689090322233579

Results for sequence length: 20

Transformer Model:
Training Loss: 2.4649602108898487
Validation Loss: 2.460609679763612
Validation Accuracy: 0.2689090322233579
Execution Time: 1632.793774843216 seconds
```

Sequences of 50:

The training loss, validation loss, validation accuracy, and total execution time can be seen in Figure 9 below:

Figure 9: Training Loss + Accuracy

```
Training Transformer model...
Epoch 1, Train Loss: 2.511785074914826, Validation Loss: 2.485132347520965, Validation Accuracy: 0.26712095400340713
Epoch 2, Train Loss: 2.482837775280943, Validation Loss: 2.477877752908963, Validation Accuracy: 0.2683807047431184
Epoch 3, Train Loss: 2.477420491778048, Validation Loss: 2.472766340489688, Validation Accuracy: 0.26829104276876176
Epoch 4, Train Loss: 2.4744677833855384, Validation Loss: 2.471256213909434, Validation Accuracy: 0.26762306105980455
Epoch 5, Train Loss: 2.4725065900425864, Validation Loss: 2.472710705158326, Validation Accuracy: 0.2633551510804268
Epoch 6, Train Loss: 2.47109892384777, Validation Loss: 2.4690889123629147, Validation Accuracy: 0.2683762216444006
Epoch 7, Train Loss: 2.4697847041657908, Validation Loss: 2.470059797969686, Validation Accuracy: 0.26894109208284767
Epoch 8, Train Loss: 2.469238147296938, Validation Loss: 2.4687171053931305, Validation Accuracy: 0.2692593920918139
Epoch 9, Train Loss: 2.4685092895840954, Validation Loss: 2.467085542450526, Validation Accuracy: 0.26848829911234645
Epoch 10, Train Loss: 2.467590626889597, Validation Loss: 2.466101540140771, Validation Accuracy: 0.26856451179054963
Epoch 11, Train Loss: 2.467185805708178, Validation Loss: 2.467029020286703, Validation Accuracy: 0.26823724558414774
Epoch 12, Train Loss: 2.466691150950026, Validation Loss: 2.4674211801781865, Validation Accuracy: 0.2689769568725903
Epoch 13, Train Loss: 2.466304288470113, Validation Loss: 2.4671611645697777, Validation Accuracy: 0.2682596610777369
Epoch 14, Train Loss: 2.466034688353592, Validation Loss: 2.466570901579999, Validation Accuracy: 0.2692907737828387
Epoch 15, Train Loss: 2.465792534719676, Validation Loss: 2.4653102789162027, Validation Accuracy: 0.269254908993096
Epoch 16, Train Loss: 2.4653771839451384, Validation Loss: 2.465447485569848, Validation Accuracy: 0.26928180758540304
Epoch 17, Train Loss: 2.4649988132624907, Validation Loss: 2.465517826930275, Validation Accuracy: 0.26895005828028334
Epoch 18, Train Loss: 2.464936914784293, Validation Loss: 2.4659160465058174, Validation Accuracy: 0.26860485967901016
Epoch 19, Train Loss: 2.464654972618881, Validation Loss: 2.4663554883479946, Validation Accuracy: 0.2708598583340805
Epoch 20, Train Loss: 2.464308036351689, Validation Loss: 2.464719433867, Validation Accuracy: 0.26953734421231956

Results for sequence length: 20

Transformer Model:
Training Loss: 2.464308036351689
Validation Loss: 2.464719433867
Validation Accuracy: 0.26953734421231956
Execution Time: 2449.5143270492554 seconds
```

Overall, the accuracy for all transformer-based models with sequences of 20, 30, and 50 were worse than the models with GRUs or LSTMs.

Problem 3.

The training loss and accuracy can be seen in Figure 10 below:

Figure 10: Training Loss + Accuracy

Epoch [1/20],	Loss: 5.9164,	Training Accuracy: 0.0160
Epoch [2/20],	Loss: 5.0685,	Training Accuracy: 0.0427
Epoch [3/20],	Loss: 4.5625,	Training Accuracy: 0.0960
Epoch [4/20],	Loss: 4.1173,	Training Accuracy: 0.1573
Epoch [5/20],	Loss: 3.7107,	Training Accuracy: 0.2427
Epoch [6/20],	Loss: 3.3208,	Training Accuracy: 0.3067
Epoch [7/20],	Loss: 2.9846,	Training Accuracy: 0.4080
Epoch [8/20],	Loss: 2.6569,	Training Accuracy: 0.4373
Epoch [9/20],	Loss: 2.3210,	Training Accuracy: 0.5040
Epoch [10/20],	Loss: 2.0154,	Training Accuracy: 0.5333
Epoch [11/20],	Loss: 1.7766,	Training Accuracy: 0.5760
Epoch [12/20],	Loss: 1.5314,	Training Accuracy: 0.6240
Epoch [13/20],	Loss: 1.2915,	Training Accuracy: 0.6853
Epoch [14/20],	Loss: 1.1415,	Training Accuracy: 0.7013
Epoch [15/20],	Loss: 0.9778,	Training Accuracy: 0.7173
Epoch [16/20],	Loss: 0.8153,	Training Accuracy: 0.7413
Epoch [17/20],	Loss: 0.6671,	Training Accuracy: 0.7680
Epoch [18/20],	Loss: 0.5718,	Training Accuracy: 0.7733
Epoch [19/20],	Loss: 0.4724,	Training Accuracy: 0.7893
Epoch [20/20],	Loss: 0.3980,	Training Accuracy: 0.7920

The validation loss and accuracy can be seen in Figure 11 below:

Figure 11: Training Loss + Accuracy

Evaluation Loss: 0.1585,	Evaluation Accuracy: 0.7947
--------------------------	-----------------------------

The training and validation loss were lower in the transformer-based model than the RNN model without attention. However, the RNN model with attention had a better training and validation loss as well as validation accuracy.

Problem 4.

The training loss and accuracy can be seen in Figure 12 below:

Figure 12: Training Loss + Accuracy

```
Epoch [1/20], Loss: 6.0092, Training Accuracy: 0.0107
Epoch [2/20], Loss: 5.1089, Training Accuracy: 0.0773
Epoch [3/20], Loss: 4.5889, Training Accuracy: 0.0987
Epoch [4/20], Loss: 4.1362, Training Accuracy: 0.1573
Epoch [5/20], Loss: 3.7471, Training Accuracy: 0.2347
Epoch [6/20], Loss: 3.3494, Training Accuracy: 0.3013
Epoch [7/20], Loss: 2.9824, Training Accuracy: 0.3493
Epoch [8/20], Loss: 2.6639, Training Accuracy: 0.4133
Epoch [9/20], Loss: 2.3408, Training Accuracy: 0.4907
Epoch [10/20], Loss: 2.0648, Training Accuracy: 0.5467
Epoch [11/20], Loss: 1.7792, Training Accuracy: 0.5893
Epoch [12/20], Loss: 1.5327, Training Accuracy: 0.6160
Epoch [13/20], Loss: 1.3426, Training Accuracy: 0.6773
Epoch [14/20], Loss: 1.1385, Training Accuracy: 0.6773
Epoch [15/20], Loss: 0.9598, Training Accuracy: 0.7280
Epoch [16/20], Loss: 0.8171, Training Accuracy: 0.7360
Epoch [17/20], Loss: 0.6916, Training Accuracy: 0.7573
Epoch [18/20], Loss: 0.5820, Training Accuracy: 0.7707
Epoch [19/20], Loss: 0.4867, Training Accuracy: 0.7760
Epoch [20/20], Loss: 0.4067, Training Accuracy: 0.7787
```

The validation loss and accuracy can be seen in Figure 13 below:

Figure 13: Training Loss + Accuracy

```
Evaluation Loss: 0.1699, Evaluation Accuracy: 0.7920
```

The RNN model with or without attention had a better training and validation loss as well as validation accuracy than the transformer-based model.